

Cloud Computing and Big Data Systems - Spring 2024 Assignment 2

Moin Khan(mk8793), Vir Jhangiani(vrj2006)

Objective:

In this assignment, you will learn how to deploy a web application on Kubernetes using Docker containers. You will start by containerizing the application using Docker, create a persistence volume, and push the image to Docker Hub. Then, you will deploy the container on a local Kubernetes cluster using Minikube and then deploy it on AWS EKS. You will also explore various Kubernetes features such as a replication controller, health monitoring, rolling updates, and alerting.

Part 1: Creating application locally

We ran the application locally to make sure it was working properly.

127.0.0.1:5000

ToDo Reminder

ALL Uncompleted Completed About

Search Reference: Unique ID Search Task Search

To-Do LIST :

Status	Task Name	Description Name	Date	Priority	Remove	Modify
X	TEST	TEST	2024-02-29	None		

Add a Task

Taskname

Enter Description here...

dd-mm-yyyy

Priority

Create

Part 2: Containerizing the Application on Docker

Wrote Dockerfile for the flask application with all required detail:

Dockerfile.yaml

```
FROM python:3.9-slim
RUN apt-get update; apt-get install iputils-ping curl -y
WORKDIR /flask-docker

COPY requirements.txt requirements.txt
RUN pip3 install -r requirements.txt

COPY . .

CMD [ "python3", "-m", "flask", "run", "--host=0.0.0.0"]
```

Docker build(Building the Docker image using the Docker CLI (command-line interface):

```
docker build . -t flask-docker-image
```

Testing the application locally, using docker-compose file that defines the services for the flask app and MongoDB container. The docker-compose file specifies the images to use, the ports to expose, and a volume to mount to store the database information.

```
docker compose up
```

docker-compose.yml

```
version: '3'
services:
  flask:
    image: flask-docker-image
    environment:
      - MONGO_HOST=mongodb
      - MONGO_PORT=27017
      - FLASK_ENV=development
      - PORT=5000
    depends_on:
      - mongodb
    ports:
      - "5000:5000"
  mongodb:
    image: mongo:4.4.15
    ports:
      - "27019:27017"
    expose:
      - "27017"
    environment:
      - MONGODB_INITDB_ROOT_USERNAME=admin
      - MONGODB_INITDB_ROOT_PASSWORD=admin
    volumes:
      - type: bind
        source: ./data
```

```
target: /data/db
```

ToDo Reminder

ALL Uncompleted Completed About

Search Reference: Unique ID Search

To Do LIST :

Status	Task Name	Description Name	Date	Priority	Remove	Modify
<input checked="" type="checkbox"/>	TODO	TEST	2024-03-18	None	[X]	[Edit]
<input checked="" type="checkbox"/>	TEST	tere	2024-03-18	None	[X]	[Edit]

Add a Task

Create

Pushing the Docker image to Docker Hub to make it available for deployment to Kubernetes.
Below are the steps using Docker CLI to push the image to the Docker Hub registry.

Docker image tag

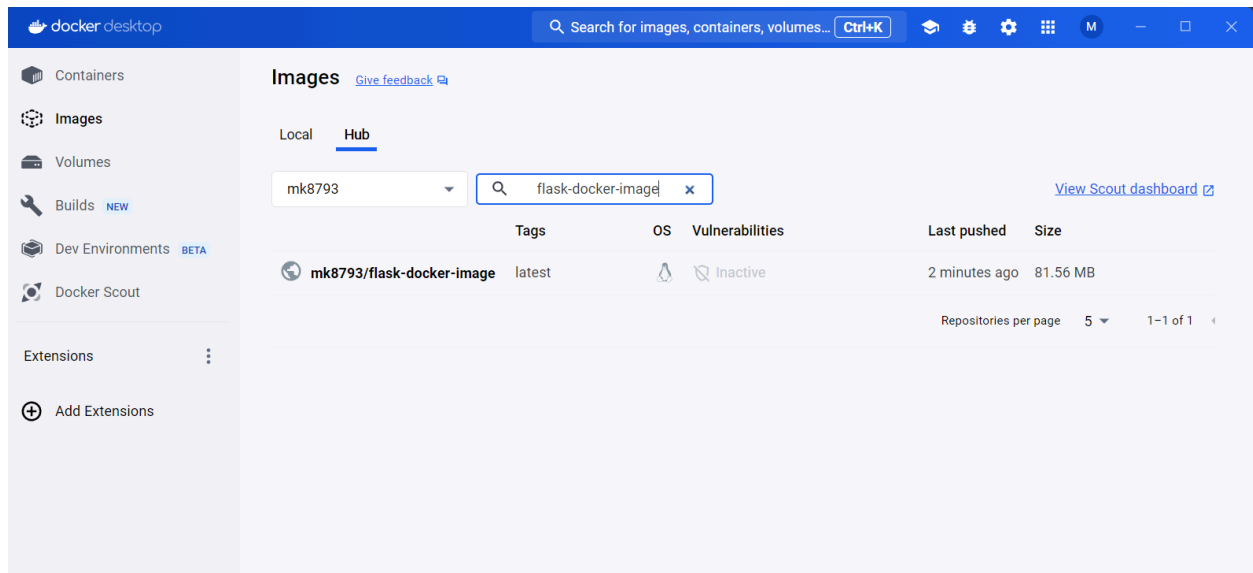
```
docker tag flask-docker-image mk8793/flask-docker-image
```

Docker push image

```
docker image push mk8793/flask-docker-image:latest
```

```
C:\Users\moink\Downloads\CC\HW2\Docker>docker image push mk8793/flask-docker-image:latest
The push refers to repository [docker.io/mk8793/flask-docker-image]
80c47523f658: Pushed
d814c5ae2bd9: Mounted from mk8793/flask-docker-new
24a0d725fa79: Mounted from mk8793/flask-docker-new
e5588c97b653: Mounted from mk8793/flask-docker-new
5c8fe12dc456: Mounted from mk8793/flask-docker-new
4a7ac3585b06: Mounted from mk8793/flask-docker-new
6be461d39d4d: Mounted from mk8793/flask-docker-new
d91aa0e727e2: Mounted from mk8793/flask-docker-new
c8f253aef560: Mounted from mk8793/flask-docker-new
a483da8ab3e9: Mounted from mk8793/flask-docker-new
latest: digest: sha256:5cd454502433a6188bd8c6b1263740c15bcd4f8f2f96b259f0eaccd65b354d31 size: 2419
```

Below you will find a screenshot of the image available on docker hub



Part 3: Deploying the Application on Minikube:

Start Minikube using the command-line interface.

```
minikube start
```

Creating PV using file mongo_pv.yml

```
kubectl apply -f mongo_pv.yml
```

```
C:\Users\moink\Downloads\CC\HW2\Minikube>kubectl apply -f mongo_pv.yml
persistentvolume/mongo-pv created

C:\Users\moink\Downloads\CC\HW2\Minikube>kubectl apply -f mongo_pv.yml

C:\Users\moink\Downloads\CC\HW2\Minikube>kubectl get pv
NAME                CAPACITY  ACCESS MODES  RECLAIM POLICY  STATUS  CLAIM                STORAGECLASS  REASON  AGE
mongo-pv            256Mi     RWO           Retain          Available                                     default/mongo-pvc  standard  17s
pvc-f221ee67-696f-4746-9847-ebd1af7246e8  256Mi     RWO           Delete          Released  default/mongo-pvc  standard  15d
```

mongo_pv.yml

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: mongo-pv
spec:
  capacity:
```

```
  storage: 256Mi
  accessModes:
    - ReadWriteOnce
  hostPath:
    path: /tmp/db
  storageClassName: "standard"
```

Create Persistent Volume Claim using mongo_pvc.yml

```
kubectl apply -f mongo_pvc.yml
```

```
C:\Users\moink\Downloads\CC\HW2\Minikube>kubectl apply -f mongo_pvc.yml
persistentvolumeclaim/mongo-pvc created
```

```
C:\Users\moink\Downloads\CC\HW2\Minikube>kubectl get pvc
```

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAGECLASS	AGE
mongo-pvc	Bound	pvc-97a17253-37d6-43ac-beb9-2c0795fd7f14	256Mi	RWO	standard	3s

mongo_pvc.yml

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: mongo-pvc
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 256Mi
```

Creating mongo Service using mongo_svc.yml:

```
Kubectl apply -f mongo_svc.yml
```

```
C:\Users\moink\Downloads\CC\HW2\Minikube>kubectl apply -f mongo_svc.yml
service/mongodb-service configured
```

```
C:\Users\moink\Downloads\CC\HW2\Minikube>kubectl get service
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	17d
mongodb-service	LoadBalancer	10.111.170.240	<pending>	27019:30218/TCP	6d3h

mongo_svc.yml

```
apiVersion: v1
kind: Service
metadata:
  name: mongodb-service
spec:
  selector:
    app: mongodb
  ports:
    - port: 27017
      targetPort: 27017
  type: LoadBalancer
```

Here the mongodb service would be accessible outside the cluster through **LoadBalancer on port 27017**. Any pod which wants to communicate with the mongodb can use the mongodb-service as a valid hostname for mongodb.

Creating mongo container in a pod using mongo_db_deployment.yml:

```
kubectl apply -f mongo_db_deployment.yml
```

```
C:\Users\moink\Downloads\CC\HW2\Minikube>kubectl apply -f mongo_db_deployment.yml
deployment.apps/mongodb-deployment created
```

mongo_db_deployment.yml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: mongodb-deployment
spec:
  replicas: 1
  selector:
    matchLabels:
```

```

    app: mongodb
template:
  metadata:
    labels:
      app: mongodb
  spec:
    containers:
      - name: mongodb-container
        image: mongodb/mongodb-community-server:6.0-ubi8
        ports:
          - containerPort: 27017
        volumeMounts:
          - name: storage
            mountPath: /data/db
        readinessProbe:
          exec:
            command:
              - /bin/sh
              - -c
              - mongo --eval "db.adminCommand('ping')"
            initialDelaySeconds: 15
            periodSeconds: 10
    volumes:
      - name: storage
        persistentVolumeClaim:
          claimName: mongo-pvc

```

The mongodb deployment yaml will create a pod with mongodb-container. The ReplicationController will manage 1 replica. To check whether mongodb db started or not, we created the readinessProbe which would check the mongodb status before sending any request to the pod.

Now deploying flask container on pod with replication controller using flask_rc.yaml:

```
kubectl apply -f flask_rc.yaml
```

```

C:\Users\moink\Downloads\CC\HW2\Minikube>kubectl apply -f flask_rc.yaml
deployment.apps/flask-rc created

```

flask_rc.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: flask-rc
spec:
  replicas: 3
  selector:
    matchLabels:
      app: flask
  template:
    metadata:
      name: flask
      labels:
        app: flask
    spec:
      containers:
        - name: flask-app-container
          image: mk8793/flask-docker-new:latest
          ports:
            - containerPort: 5000
          env:
            - name: MONGO_HOST
              value: mongodb-service
            - name: MONGO_PORT
              value: "27017"
            - name: MONGODB_USERNAME
              value: admin
            - name: MONGODB_PASSWORD
              value: admin
            - name: FLASK_ENV
              value: development
            - name: PORT
              value: "5000"
      livenessProbe:
        httpGet:
          path: /
          port: 5000
        initialDelaySeconds: 10
```



```

        periodSeconds: 30
        failureThreshold: 3
    readinessProbe:
        exec:
            command:
            - /bin/sh
            - -c
            - curl --fail http://mongodb-service:27017/ || exit 1
    initialDelaySeconds: 10
    periodSeconds: 30
    failureThreshold: 4

```

To test the app locally, run

```
minikube tunnel
```

ToDo Reminder

ALL
Uncompleted
Completed
About

Search Reference:

Status	Task Name	Description Name	Date	Priority	Remove	Modify
✖	TODO Minikube	Minikube	2024-03-18	None		

Add a Task

Part 5: Replication controller feature:

In the YAML file above flask deployment will pull the image from mk8793/flask-docker-image:latest from the docker hub. The deployment also has replicationController enabled, which would manage 3 replicas, i.e 3 pods should be there at least incase of failure/scale down.

To Observe our running pods including the Replication Controller(3 pods for flask):

```
minikube dashboard
```

Name	Images	Labels	Node	Status	Restarts	CPU Usage (cores)	Memory Usage (bytes)	Created
flask-rc-6f78cf8dfd-w4mbp	mk8793/flask-docker-new:latest	app: flask pod-template-hash: 6f78cf8dfd	minikube	Running	0	-	-	17 minutes ago
flask-rc-6f78cf8dfd-tgt78	mk8793/flask-docker-new:latest	app: flask pod-template-hash: 6f78cf8dfd	minikube	Running	0	-	-	17 minutes ago
flask-rc-6f78cf8dfd-f2wv9	mk8793/flask-docker-new:latest	app: flask pod-template-hash: 6f78cf8dfd	minikube	Running	0	-	-	17 minutes ago
mongodb-deployment-7769d6dbd-dmczg	mongodb/mongodb-community-server:6.0-ubi8	app: mongodb pod-template-hash: 7769d6dbd	minikube	Running	0	-	-	17 minutes ago

Part 7: Health monitoring

The deployment has two probes you can observe their definition in the flask_rc.yml file above:

- **Liveness probe**, which sends an HTTP Get request to the flask application to check whether the flask is active or not.
- **Readiness Probe** checks whether the flask application is able to make a connection with the mongodb container before accepting any request.

Part 6: RollingUpdate

Update docker image version to v2

```
C:\Users\moink\Downloads\CC\HW2\TODO>docker tag flask-docker-image mk8793/flask-docker-image:v2

C:\Users\moink\Downloads\CC\HW2\TODO>docker image push mk8793/flask-docker-image:v2
The push refers to repository [docker.io/mk8793/flask-docker-image]
80c47523f658: Layer already exists
d814c5ae2bd9: Layer already exists
24a0d725fa79: Layer already exists
e5588c97b653: Layer already exists
5c8fe12dc456: Layer already exists
4a7ac3585b06: Layer already exists
6be461d39d4d: Layer already exists
d91aa0e727e2: Layer already exists
c8f253aef560: Layer already exists
a483da8ab3e9: Layer already exists
v2: digest: sha256:5cd454502433a6188bd8c6b1263740c15bcd4f8f2f96b259f0eaccd65b354d31 size: 2419
```

Then update kubernetes image from latest to v2

```
C:\Users\moink\Downloads\CC\HW2\TODO>kubectl set image deployments/flask-rc flask-app-container=mk8793/flask-docker-image:v2
deployment.apps/flask-rc image updated
```

Below you can observe that our image has changed from

mk8793/flask-docker-new:latest

to

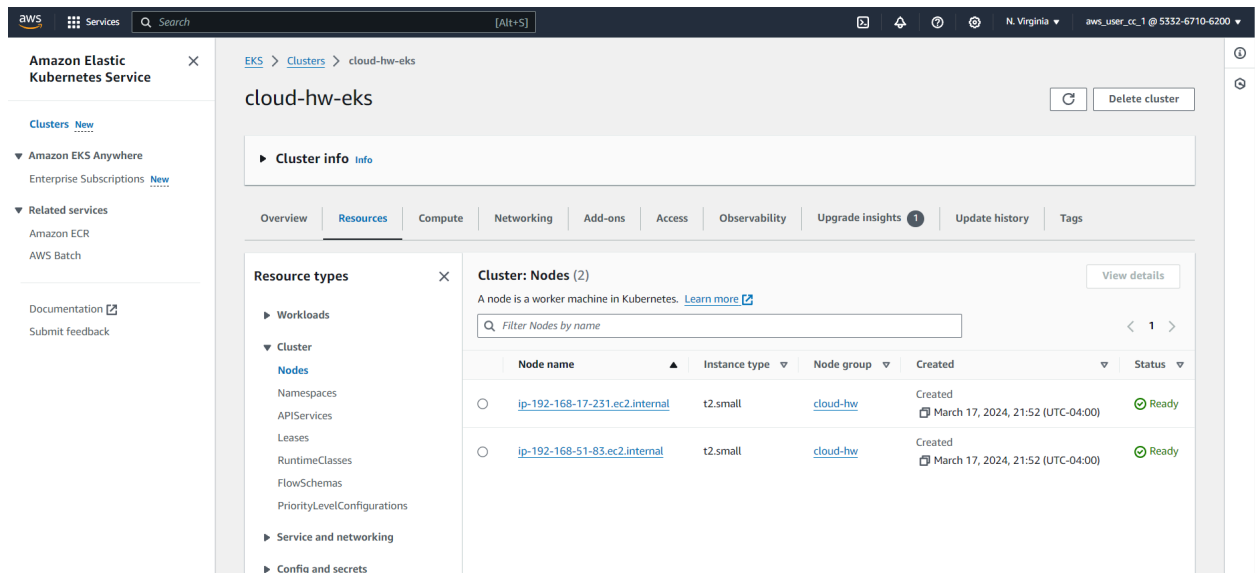
mk8793/flask-docker-new:v2

Name	Images	Labels	Node	Status	Restarts	CPU Usage (cores)	Memory Usage (bytes)	Created
flask-rc-8dcc8f46d-cqhjrc	mk8793/flask-docker-image:v2	app: flask pod-template-hash: 8dcc8f46d	minikube	Running	0	-	-	23 seconds ago
flask-rc-8dcc8f46d-fhjrt	mk8793/flask-docker-image:v2	app: flask pod-template-hash: 8dcc8f46d	minikube	Running	0	-	-	27 seconds ago
flask-rc-8dcc8f46d-lb4sv	mk8793/flask-docker-image:v2	app: flask pod-template-hash: 8dcc8f46d	minikube	Running	0	-	-	35 seconds ago
mongodb-deployment-7769d6dbd-dmczg	mongodb/mongodb-community-server.6.0-ubi8	app: mongodb pod-template-hash: 7769d6dbd	minikube	Running	0	-	-	24 minutes ago

Part 4: Deploying the Application on AWS EKS

Create cluster:

```
eksctl create cluster --name cloud-hw-eks --region us-east-1
--nodegroup-name cloud-hw --node-type t2.small --nodes 3 --nodes-min 1
--nodes-max 5 --managed
```



IAM permission for IAM role:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "VisualEditor0",
      "Effect": "Allow",
      "Action": [
        "iam:Attach*",
        "autoscaling:Describe*",
        "iam:List*",
        "ec2:DeleteTags",
        "iam:UntagRole",
        "iam:TagRole",
        "iam:TagMFADevice",
        "iam:TagSAMLProvider",
        "iam:Put*",
        "iam:PassRole",
        "iam:Get*",
        "autoscaling:UpdateAutoScalingGroup",
        "iam:UntagSAMLProvider",

```

```

        "ssm:Get*",
        "ec2:CreateTags",
        "iam:Create*",
        "iam:Detach*",
        "cloudformation:*",
        "iam:UntagServerCertificate",
        "iam:TagPolicy",
        "iam:TagUser",
        "ec2:CreateVolume",
        "iam:UntagUser",
        "iam:Delete*",
        "ec2:Describe*",
        "iam:Update*",
        "iam:UntagMFADevice",
        "iam:TagServerCertificate",
        "iam:UntagPolicy",
        "ssm:List*",
        "iam:UntagOpenIDConnectProvider",
        "eks:*",
        "iam:UntagInstanceProfile",
        "iam:TagOpenIDConnectProvider",
        "iam:TagInstanceProfile"
    ],
    "Resource": "*"
}
]
}

```

Configure Kubernetes:

```
aws eks update-kubeconfig --region us-east-1 --name cloud-hw-eks
```

Create OIDC Provider:

```
eksctl utils associate-iam-oidc-provider --cluster cloud-hw-eks --approve
```

Check OIDC:

```
aws iam list-open-id-connect-providers
```

Create Storage Class (EBS) for persistentVolumes (required for MongoDB) using
kubectl create -f eks_ebs_sc.yml

```
C:\Users\moink\Downloads\CC\HW2\EKS>kubectl create -f eks_ebs_sc.yml
storageclass.storage.k8s.io/gp2 created

C:\Users\moink\Downloads\CC\HW2\EKS>kubectl get storageclass
NAME                PROVISIONER          RECLAIMPOLICY   VOLUMEBINDINGMODE   ALLOWVOLUMEEXPANSION   AGE
gp2 (default)       kubernetes.io/aws-ebs   Delete          WaitForFirstConsumer   false                  6s
```

eks_ebs_sc.yml

```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  annotations:
    storageclass.kubernetes.io/is-default-class: "true"
  name: gp2
parameters:
  fsType: ext4
  type: gp2
provisioner: kubernetes.io/aws-ebs
volumeBindingMode: WaitForFirstConsumer
```

Create IAM Role:

```
eksctl create iamserviceaccount \
  --name ebs-csi-controller-sa \
  --namespace kube-system \
  --cluster cloud-hw-eks \
  --role-name AmazonEKS_EBS_CSI_DriverRole \
  --role-only \
  --attach-policy-arn
arn:aws:iam::aws:policy/service-role/AmazonEBSCSIDriverPolicy \
  --approve
```

Attach policy:

```
aws iam attach-role-policy --policy-arn
arn:aws:iam::533267106200:role/AmazonEKS_EBS_CSI_DriverRole --role-name
AmazonEKS_EBS_CSI_DriverRole
```

Create Amazon EBS CSI Driver Addon:

```
eksctl create addon --name aws-ebs-csi-driver --cluster cloud-hw-eks
--service-account-role-arn
arn:aws:iam::533267106200:role/AmazonEKS_EBS_CSI_DriverRole --force
```

Create PersistentVolumeClaim (Do not require PersistentVolume in Dynamic Provisioning)

kubectl create -f eks_ekspvc.yml

eks_ekspvc.yml

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: mongo-pvc
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 256Mi
  storageClassName: gp2
```

MongoDb_EKS_Deployment.yml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: mongodb-deployment
spec:
  replicas: 1
  selector:
    matchLabels:
```

```

    app: mongodb
template:
  metadata:
    labels:
      app: mongodb
  spec:
    containers:
      - name: mongodb-container
        image: mongo:4.4.15
        ports:
          - containerPort: 27017
        env:
          - name: MONGODB_INITDB_ROOT_USERNAME
            value: admin
          - name: MONGODB_INITDB_ROOT_PASSWORD
            value: admin
        volumeMounts:
          - name: storage
            mountPath: /data/db
        readinessProbe:
          exec:
            command:
              - /bin/sh
              - -c
              - mongo --eval "db.adminCommand('ping')"
            initialDelaySeconds: 15
            periodSeconds: 10
    volumes:
      - name: storage
        persistentVolumeClaim:
          claimName: mongo-pvc

```

Creating mongo service on AWS EKS:

```

apiVersion: v1
kind: Service
metadata:
  name: mongodb-service
spec:
  selector:

```



```
app: mongodb
ports:
  - port: 27017
    targetPort: 27017
type: LoadBalancer
```

Status of mongo deployment on EKS:

The screenshot shows the AWS Management Console interface for the Amazon Elastic Kubernetes Service. The breadcrumb navigation indicates the path: EKS > Clusters > cloud-hw-eks > Pod: default/mongodb-deployment-85c496c6c-4r2bm. The main heading is 'mongodb-deployment-85c496c6c-4r2bm'. The 'Info' section displays the following details:

Info		
Status: Running	Created: 3 minutes ago	Namespace: default
Controlled by: ReplicaSet/mongodb-deployment-85c496c6c	Last transition time: 3 minutes ago	Node: ip-192-168-51-83.ec2.internal
Volume: Volumes:tag.kubernetes.io/created-for/pvc/name=mongo-pvc	Pod IP: 192.168.57.132	

The 'Containers (1)' section shows a single container named 'mongodb-container'. The 'Labels (2)' and 'Annotations (0)' sections are currently empty.

Mongo service on EKS:

The screenshot shows the AWS Management Console interface for the Amazon Elastic Kubernetes Service, specifically the configuration page for a service. The breadcrumb navigation indicates the path: EKS > Clusters > cloud-hw-eks > Service: default/mongodb. The main heading is 'mongodb'. The 'Info' section displays the following details:

Info		
Created: a few seconds ago	Namespace: default	Selector: app=mongodb
Finalizers: service.kubernetes.io/load-balancer-cleanup	Type: LoadBalancer	Session affinity: None
IP family policy: SingleStack	Cluster IPs: 10.100.136.81	IP families: IPv4
Load balancer URLs: acc9c87666246439780c112cfc72a0fb-417878328.us-east-1.elb.amazonaws.com		

The 'Ports (1)' section shows a single port configuration:

Protocol	Port	Target port	Node port
TCP	27017	27017	30600

The 'Endpoints (1)' section shows a single endpoint configuration:

IP	Hostname	Node name	Target reference
192.168.57.132:27017	-	ip-192-168-51-83.ec2.internal	mongodb-deployment-85c496c6c-4r2bm

Observe the running mongo deployment and that mongo is accessed outside the cluster:

```
C:\Users\moink>curl http://acc9c87666246439780c112cfc72a0fb-417878328.us-east-1.elb.amazonaws.com:27017
It looks like you are trying to access MongoDB over HTTP on the native driver port.
```

```
C:\Users\moink\Downloads\CC\HW2\EKS>kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
mongodb-deployment-85c496c6c-4r2bm  1/1     Running   0           82s
```

Flask deployment on EKS:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: flask-rc
spec:
  replicas: 3
  selector:
    matchLabels:
      app: flask
  template:
    metadata:
      name: flask
      labels:
        app: flask
    spec:
      containers:
      - name: flask-app-container
        image: mk8793/flask-docker-new:latest
        ports:
        - containerPort: 5000
        env:
          - name: MONGO_HOST
            value: mongodb-service
          - name: MONGO_PORT
            value: "27017"
          - name: FLASK_ENV
            value: development
          - name: PORT
            value: "5000"
        livenessProbe:
          httpGet:
```

```
    path: /
    port: 5000
    initialDelaySeconds: 10
    periodSeconds: 30
    failureThreshold: 3
  readinessProbe:
    exec:
      command:
        - /bin/sh
        - -c
        - curl --fail http://mongodb-service:27017/ || exit 1
    initialDelaySeconds: 10
    periodSeconds: 30
    failureThreshold: 4
```


Creating Service for flask on EKS:

```
apiVersion: v1
kind: Service
metadata:
  name: flask-svc-lb
spec:
  selector:
    app: flask
  ports:
    - port: 5000
      targetPort: 5000
  type: LoadBalancer
```

Check Pods on EKS - 3 replicas for flask and one pod for mongodb:

EKS > Clusters > cloud-hw-eks

cloud-hw-eks

 Delete cluster

Cluster info Info

Status Active	Kubernetes version 1.29	Support type Standard support until March 23, 2025	Provider EKS
------------------	----------------------------	---	-----------------

OverviewResourcesComputeNetworkingAdd-onsAccessObservabilityUpgrade insights1Update historyTags

Resource types ×

Workloads

- PodTemplates
- Pods**
- ReplicaSets
- Deployments
- StatefulSets
- DaemonSets
- Jobs
- CronJobs
- PriorityClasses
- HorizontalPodAutoscalers

Workloads: Pods (4)

Pod is the smallest and simplest Kubernetes object. A Pod represents a set of running containers on your cluster. [Learn more](#)

default

	Name	Age
<input type="radio"/>	flask-rc-7bb87847cb-6lm78	Created a minute ago
<input type="radio"/>	flask-rc-7bb87847cb-bbhpr	Created a minute ago
<input type="radio"/>	flask-rc-7bb87847cb-p7w98	Created a minute ago
<input type="radio"/>	mongodb-deployment-85c496c6c-4r2bm	Created 8 minutes ago


View details

< 1 >

Observe the Flask Service on EKS and access it using load balancer url:

EKS > Clusters > cloud-hw-eks > Service: default/flask-svc-lb

flask-svc-lb

Structured viewRaw view

Info

Created 2 minutes ago	Namespace default	Selector app=flask
Finalizers service.kubernetes.io/load-balancer-cleanup	Type LoadBalancer	Session affinity None
IP family policy SingleStack	Cluster IPs 10.100.229.35	IP families IPv4
Load balancer URLs ae7a37c73e4a34e479d4fba7af9ae3bd-165991769.us-east-1.elb.amazonaws.com		

Ports (1)

Protocol	Port	Target port	Node port
TCP	5000	5000	31242

Labels (0)

Annotations (1)

< 1 >

< 1 >

Access flask from outside:

← → ⚠ Not secure ae7a37c73e4a34e479d4fba7af9ae3bd-165991769.us-east-1.elb.amazonaws.com:5000/list

ToDo Reminder

ALL Uncompleted Completed About

Search Reference: Unique ID Search Task Search

To-Do LIST :

Status	Task Name	Description Name	Date	Priority	Remove	Modify
X	AWS TASK	TASK	2024-03-18	None		

Add a Task

Taskname

Enter Description here...

dd-mm-yyyy

Priority

Create

Step 8: Alerting (Extra Credit 20 Points):

Deploy prometheus

Create namespace for prometheus:

```
kubectl create namespace monitoring
```

Create Cluster Role:

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: prometheus
rules:
- apiGroups: [""]
  resources:
    - nodes
    - nodes/proxy
    - services
    - endpoints
    - pods
  verbs: ["get", "list", "watch"]
- apiGroups:
  - extensions
  resources:
    - ingresses
  verbs: ["get", "list", "watch"]
```

```

- nonResourceURLs: ["/metrics"]
  verbs: ["get"]
---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: prometheus
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: prometheus
subjects:
- kind: ServiceAccount
  name: default
  namespace: monitoring

```

Create config-map:

```

apiVersion: v1
kind: ConfigMap
metadata:
  name: prometheus-server-conf
  labels:
    name: prometheus-server-conf
  namespace: monitoring
data:
  prometheus.rules: |-
    groups:
    - name: prometheus_alert
      rules:
      - alert: High Pod Memory
        expr: sum(container_memory_usage_bytes) > 1
        for: 1m
        labels:
          severity: slack
        annotations:

```

```

        summary: High Memory Usage
    - name: failing pods
      rules:
        - alert: PodFailingProbes
          expr:
sum(kube_pod_container_status_waiting_reason{reason=~"ProbeFailed|CrashLoopBackOff|ImagePullBackOff|InvalidImageName"}) by (namespace, pod) > 0
          for: 1m
          labels:
            severity: slack
          annotations:
            summary: Failing probs
    - name: pod-restart-alerts
      rules:
        - alert: PodRestartExceedsThreshold
          expr: increase(kube_pod_container_status_restarts_total[10m]) >
3
          for: 1m
          labels:
            severity: slack
          annotations:
            summary: "Pod {{ $labels.pod }} in namespace {{
$labels.namespace }} has restarted more than 3 times"
            description: "The pod {{ $labels.pod }} in namespace {{
$labels.namespace }} has restarted more than 3 times within the last 5
minutes."
prometheus.yml: |-
  global:
    scrape_interval: 5s
    evaluation_interval: 5s
  rule_files:
    - /etc/prometheus/prometheus.rules
  alerting:
    alertmanagers:
      - scheme: http
        static_configs:
          - targets:
            - "alertmanager.monitoring.svc:9093"
  scrape_configs:
    - job_name: 'node-exporter'

```

```
kubernetes_sd_configs:
  - role: endpoints
relabel_configs:
- source_labels: [__meta_kubernetes_endpoints_name]
  regex: 'node-exporter'
  action: keep
- job_name: 'kubernetes-apiservers'
  kubernetes_sd_configs:
    - role: endpoints
  scheme: https
  tls_config:
    ca_file: /var/run/secrets/kubernetes.io/serviceaccount/ca.crt
  bearer_token_file:
/var/run/secrets/kubernetes.io/serviceaccount/token
  relabel_configs:
    - source_labels: [__meta_kubernetes_namespace,
__meta_kubernetes_service_name, __meta_kubernetes_endpoint_port_name]
      action: keep
      regex: default;kubernetes;https
- job_name: 'kubernetes-nodes'
  scheme: https
  tls_config:
    ca_file: /var/run/secrets/kubernetes.io/serviceaccount/ca.crt
  bearer_token_file:
/var/run/secrets/kubernetes.io/serviceaccount/token
  kubernetes_sd_configs:
    - role: node
  relabel_configs:
    - action: labelmap
      regex: __meta_kubernetes_node_label_(.+)
    - target_label: __address__
      replacement: kubernetes.default.svc:443
    - source_labels: [__meta_kubernetes_node_name]
      regex: (.+)
      target_label: __metrics_path__
      replacement: /api/v1/nodes/${1}/proxy/metrics
- job_name: 'kubernetes-pods'
  kubernetes_sd_configs:
    - role: pod
  relabel_configs:
```



```
- source_labels:
[__meta_kubernetes_pod_annotation_prometheus_io_scrape]
  action: keep
  regex: true
- source_labels:
[__meta_kubernetes_pod_annotation_prometheus_io_path]
  action: replace
  target_label: __metrics_path__
  regex: (.+)
- source_labels: [__address__,
__meta_kubernetes_pod_annotation_prometheus_io_port]
  action: replace
  regex: ([^:]+)(?::\d+)?;(\d+)
  replacement: $1:$2
  target_label: __address__
- action: labelmap
  regex: __meta_kubernetes_pod_label_(.+)
- source_labels: [__meta_kubernetes_namespace]
  action: replace
  target_label: kubernetes_namespace
- source_labels: [__meta_kubernetes_pod_name]
  action: replace
  target_label: kubernetes_pod_name
- job_name: 'kube-state-metrics'
  static_configs:
    - targets:
['kube-state-metrics.kube-system.svc.cluster.local:8080']
- job_name: 'kubernetes-cadvisor'
  scheme: https
  tls_config:
    ca_file: /var/run/secrets/kubernetes.io/serviceaccount/ca.crt
    bearer_token_file:
/var/run/secrets/kubernetes.io/serviceaccount/token
  kubernetes_sd_configs:
    - role: node
  relabel_configs:
    - action: labelmap
      regex: __meta_kubernetes_node_label_(.+)
    - target_label: __address__
      replacement: kubernetes.default.svc:443
```

```

- source_labels: [__meta_kubernetes_node_name]
  regex: (.+)
  target_label: __metrics_path__
  replacement: /api/v1/nodes/${1}/proxy/metrics/cadvisor
- job_name: 'kubernetes-service-endpoints'
  kubernetes_sd_configs:
  - role: endpoints
  relabel_configs:
  - source_labels:
    [__meta_kubernetes_service_annotation_prometheus_io_scrape]
    action: keep
    regex: true
  - source_labels:
    [__meta_kubernetes_service_annotation_prometheus_io_scheme]
    action: replace
    target_label: __scheme__
    regex: (https?)
  - source_labels:
    [__meta_kubernetes_service_annotation_prometheus_io_path]
    action: replace
    target_label: __metrics_path__
    regex: (.+)
  - source_labels: [__address__,
    __meta_kubernetes_service_annotation_prometheus_io_port]
    action: replace
    target_label: __address__
    regex: ([^:]+)(?::\d+)?;(\d+)
    replacement: $1:$2
  - action: labelmap
    regex: __meta_kubernetes_service_label_(.+)
  - source_labels: [__meta_kubernetes_namespace]
    action: replace
    target_label: kubernetes_namespace
  - source_labels: [__meta_kubernetes_service_name]
    action: replace
    target_label: kubernetes_name

```

Prometheus-deployment.yaml

apiVersion: apps/v1

```
kind: Deployment
metadata:
  name: prometheus-deployment
  namespace: monitoring
  labels:
    app: prometheus-server
spec:
  replicas: 1
  selector:
    matchLabels:
      app: prometheus-server
  template:
    metadata:
      labels:
        app: prometheus-server
    spec:
      containers:
        - name: prometheus
          image: prom/prometheus
          args:
            - "--config.file=/etc/prometheus/prometheus.yml"
            - "--storage.tsdb.path=/prometheus/"
          ports:
            - containerPort: 9090
          volumeMounts:
            - name: prometheus-config-volume
              mountPath: /etc/prometheus/
            - name: prometheus-storage-volume
              mountPath: /prometheus/
      volumes:
        - name: prometheus-config-volume
          configMap:
            defaultMode: 420
            name: prometheus-server-conf
        - name: prometheus-storage-volume
          emptyDir: {}
```

Create prometheus deployment:

```
kubect1 create -f prometheus-deployment.yaml
```

prometheus_service.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: prometheus-service
  namespace: monitoring
  annotations:
    prometheus.io/scrape: 'true'
    prometheus.io/port: '9090'
spec:
  selector:
    app: prometheus-server
  type: NodePort
  ports:
    - port: 8080
      targetPort: 9090
      nodePort: 30000
```

Creating prometheus service:

```
kubect1 create -f prometheus-service.yaml --namespace=monitoring
```

Run the following to access prometheus-service

```
minikube service prometheus-service --namespace=monitoring
```

Rules for alerts:

1. High Pod Memory Alert:
Alert Name: High Pod Memory

Expression (expr): `sum(container_memory_usage_bytes) > 1`

Condition: If the sum of container memory usage bytes exceeds 1 for 1 minute.

Labels: Sets severity label to "slack".

Annotations: Provides a summary message "High Memory Usage" for the alert.

2. Failing Pods Alert:

Alert Name: PodFailingProbes

Expression (expr):

`sum(kube_pod_container_status_waiting_reason{reason=~"ProbeFailed|CrashLoopBackOff|ImagePullBackOff|InvalidImageName"}) by (namespace, pod) > 0`

Condition: If there is at least one pod in a failing state (ProbeFailed, CrashLoopBackOff, ImagePullBackOff, InvalidImageName) within a namespace.

Labels: Sets severity label to "slack".

Annotations: Provides a summary message "Failing probes" for the alert.

3. Pod Restart Alerts:

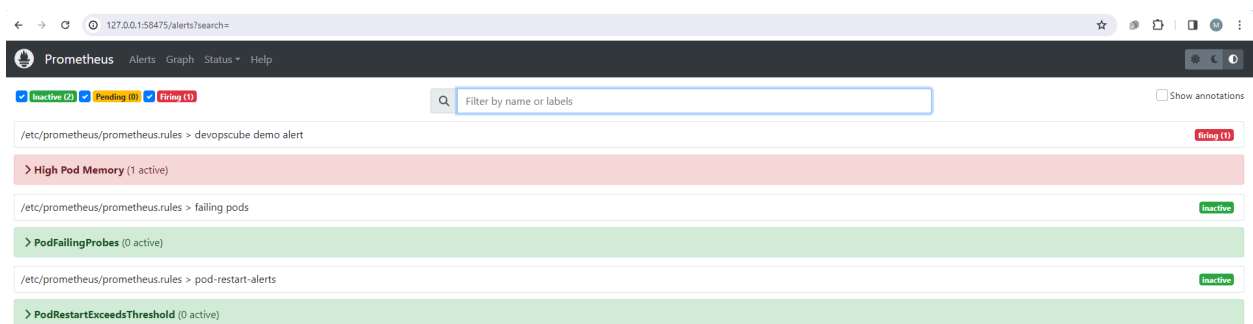
Alert Name: PodRestartExceedsThreshold

Expression (expr): `increase(kube_pod_container_status_restarts_total[10m]) > 3`

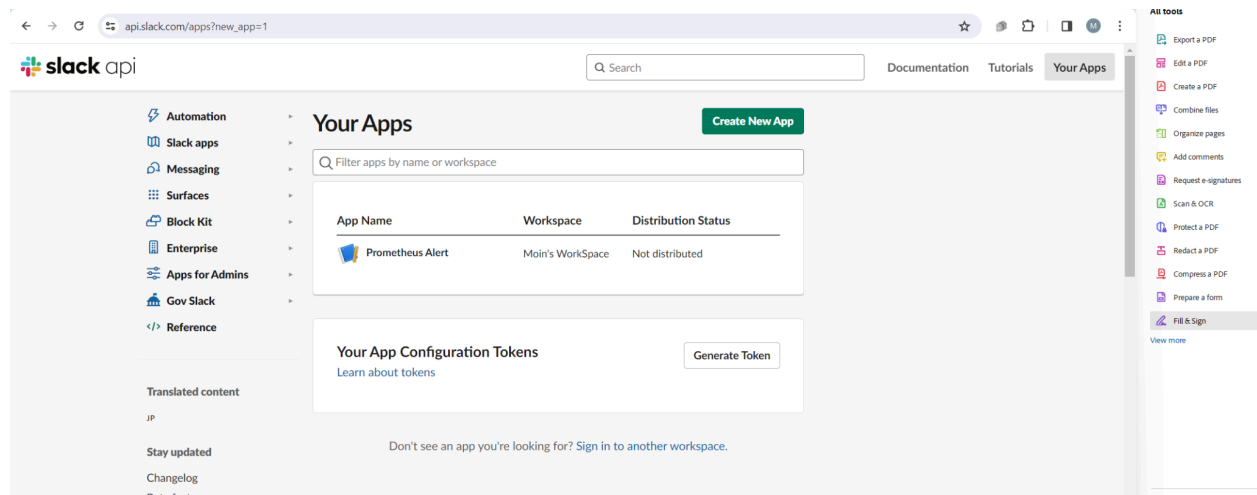
Condition: If the number of pod restarts exceeds 3 within the last 10 minutes.

Labels: Sets severity label to "slack".

Annotations: Provides a summary message "Pod {{ \$labels.pod }} in namespace {{ \$labels.namespace }} has restarted more than 3 times" and a description message "The pod {{ \$labels.pod }} in namespace {{ \$labels.namespace }} has restarted more than 3 times within the last 5 minutes."



Create Slack APP Webhook:



Alert Manager:

Add config of alertmanager to Prometheus config map

(we have already added in above Prometheus Config Map)

```
alerting:
  alertmanagers:
    - scheme: http
      static_configs:
        - targets:
          - "alertmanager.monitoring.svc:9093"
```

Crate config for alert manager:

```
kind: ConfigMap
apiVersion: v1
metadata:
  name: alertmanager-config
  namespace: monitoring
data:
  config.yml: |-
    global:
    templates:
    - '/etc/alertmanager/*.tmpl'
```

```

route:
  receiver: slack_demo
  group_by: ['alertname', 'priority']
  group_wait: 10s
  repeat_interval: 30m
  routes:
    - receiver: slack_demo
      # Send severity=slack alerts to slack.
      match:
        severity: slack
      group_wait: 10s
      repeat_interval: 1m

receivers:
- name: slack_demo
  slack_configs:
    - api_url:
https://hooks.slack.com/services/T06PJA3QQUC/B06P98KGUA2/5EoPKYkvNen3m9ObL
SHt15Si
      channel: '#prometheus_alerts'

```

Here slack_configs stores has the configuration of slack service api url along with channel where the alerts will be posted

Create alert manager deployment:

alert_manger_deployment.yaml

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: alertmanager
  namespace: monitoring
spec:
  replicas: 1
  selector:
    matchLabels:
      app: alertmanager
  template:
    metadata:
      name: alertmanager
      labels:

```

```

    app: alertmanager
spec:
  containers:
  - name: alertmanager
    image: prom/alertmanager:latest
    args:
      - "--config.file=/etc/alertmanager/config.yml"
      - "--storage.path=/alertmanager"
    ports:
      - name: alertmanager
        containerPort: 9093
    resources:
      requests:
        cpu: 500m
        memory: 500M
      limits:
        cpu: 1
        memory: 1Gi
    volumeMounts:
      - name: config-volume
        mountPath: /etc/alertmanager
      - name: templates-volume
        mountPath: /etc/alertmanager-templates
      - name: alertmanager
        mountPath: /alertmanager
  volumes:
  - name: config-volume
    configMap:
      name: alertmanager-config
  - name: templates-volume
    configMap:
      name: alertmanager-templates
  - name: alertmanager
    emptyDir: {}

```

Create Service to access the alert manager:

```

apiVersion: v1
kind: Service

```



```

metadata:
  name: alertmanager
  namespace: monitoring
  annotations:
    prometheus.io/scrape: 'true'
    prometheus.io/port: '9093'
spec:
  selector:
    app: alertmanager
  type: NodePort
  ports:
    - port: 9093
      targetPort: 9093
      nodePort: 31000

```

Alert manager Pods can be observed below on our dashboard:

The screenshot shows the Kubernetes dashboard interface. The top navigation bar includes the Kubernetes logo, a search bar, and a dropdown menu set to 'monitoring'. The main content area is titled 'Workloads > Pods'. On the left sidebar, the 'Pods' option is selected under the 'Workloads' section. The main table displays two pods:

Name	Images	Labels	Node	Status	Restarts	CPU Usage (cores)	Memory Usage (bytes)	Created
alertmanager-6c549d59d-dkgzk	prom/alertmanager:latest	app: alertmanager pod-template-hash: 6c549d59d	minikube	Running	2	-	-	4 days ago
prometheus-deployment-6d6ff75fbb-76h5s	prom/prometheus	app: prometheus-server pod-template-hash: 6d6ff75fbb	minikube	Running	2	-	-	4 days ago

Alert Manager Dashboard:

